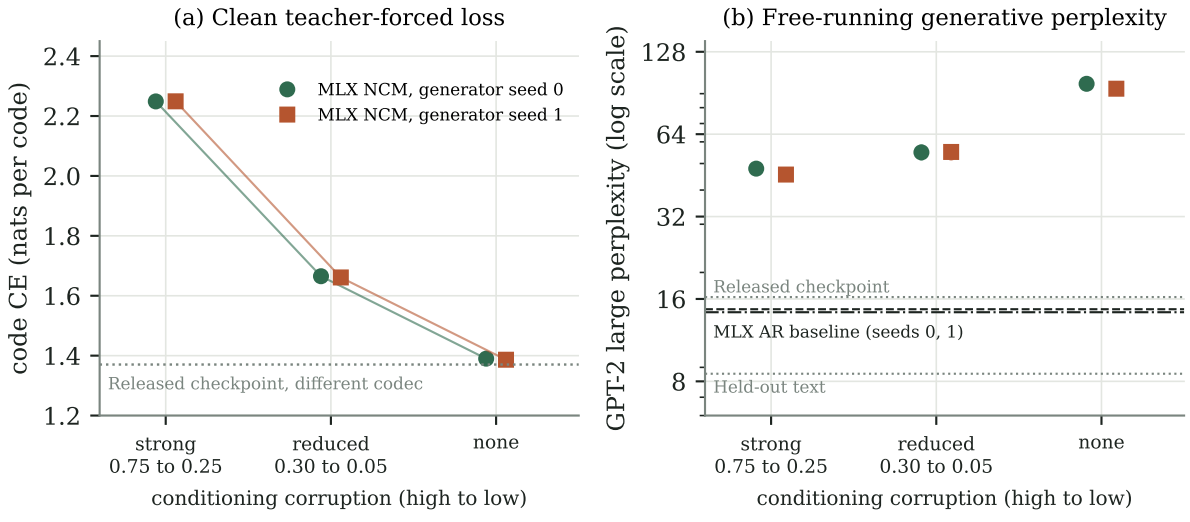
(a) Clean teacher-forced loss
code CE (nats per code)
2.4
2.2
2.0
1.8
1.6
1.4
1.2
MLX NCM, generator seed 0
MLX NCM, generator seed 1
Released checkpoint, different codec
strong
0.75 to 0.25
reduced
0.30 to 0.05
none
conditioning corruption (high to low)
(b) Free-running generative perplexity
GPT-2 large perplexity (log scale)
128
64
32
16
8
Released checkpoint
MLX AR baseline (seeds 0, 1)
Held-out text
strong
0.75 to 0.25
reduced
0.30 to 0.05
none
conditioning corruption (high to low)